

# **Protein 2D Structure Prediction Using Deep Learning**

**Name:** Md. Ibne Sina

**Roll:** 1907002

**Year:** 4th

**Term:** 2nd

**Date:** 17 August 2025

# 1. Introduction

Protein secondary structure prediction maps an amino-acid sequence to per-residue labels such as helix, sheet, and coil. Accurate 2D (secondary) structure is a key intermediate toward 3D modeling and is widely used in fold recognition, functional annotation, and downstream tasks like contact prediction. This project implements a compact convolutional neural network that performs per-residue classification directly from sequence one-hot encodings and evaluates it with standard metrics and visualizations. In addition to model evaluation, training dynamics are visualized by generating GIFs of learned convolutional filters across epochs.

## 2. Data

**Sources and format.** Sequences are provided in FASTA format (`sequence.fasta`). For each FASTA record `>ID`, a corresponding secondary-structure file `Protein Structures/ID.ss2` is parsed to obtain per-residue labels. The `.ss2` format follows the common PSIPRED output layout: each data line contains the residue index, amino acid letter, and predicted secondary-structure label.

**Label mapping.** Labels are mapped to the Q3 set used by this project:

- Helix: `H` (includes any helix subtype consolidated as `H` in parse)
- Sheet: `B` (the code maps `E`  $\rightarrow$  `B` to unify all  $\beta$ -strand states)
- Coil/Other: `C`

**Tokenization and padding.** Each sequence is one-hot encoded over the 20 canonical amino acids `ACDEFGHIKLMNPQRSTVWY`. Sequences are padded with zeros to the maximum length in the batch (`max_len`). Labels are padded with the majority class `C` to align shapes, and masked out during evaluation plots and metrics to avoid padding bias.

**Splits.** The dataset is split into training, validation, and test sets using a 70/15/15 split with `random_state=42`.

## 3. Methods

### 3.1 Preprocessing

- **One-hot encoding:** ( $L \times 20$ ) tensor per sequence, where  $L$  is padded length.
- **Label encoding:** per-residue labels are mapped to integers `{H:0, C:1, B:2}` and converted to one-hot ( $L \times 3$ ) for categorical cross-entropy.

- **Masking in analysis:** During confusion-matrix and report computation, positions whose input rows are all zeros are filtered out so evaluation is computed only on real residues.

### 3.2 Model architecture

A simple yet strong **fully convolutional** network performs per-position classification. The final layer is a  $1 \times 1$  convolution producing a 3-way distribution at each residue.

```
Input: sequence one-hot (L × 20)
└─ Conv1D(64, k=3, padding='same') + ReLU
    └─ Dropout(0.3)
        └─ Conv1D(128, k=3, padding='same') + ReLU
            └─ Dropout(0.3)
                └─ Conv1D(3, k=1, padding='same') + Softmax → (L × 3)

Loss: Categorical Cross-Entropy (per residue)
Optimizer: Adam (default settings)
Metrics: Accuracy (Keras), plus detailed metrics computed post-hoc with masking
```

This architecture captures local motifs with small receptive fields while remaining length-preserving, enabling per-residue outputs aligned with the input.

### 3.3 Training and logging

- **Epochs:** 100
- **Loss:** categorical\_crossentropy
- **Optimizer:** Adam
- **Callbacks:** A custom `WeightsGIFCallback` records the first convolutional layer's weight tensor at the end of each epoch. A heatmap frame is saved per epoch and combined into a GIF at the end of training.

**Weight-evolution visualization:** `weights_frames/weights_evolution.gif`

### 3.4 Training curves




Fig – 1: Loss over epoch




Fig – 2: Accuracy over epoch

### 3.5 Confusion matrix




Fig – 3: Confusion Matrix

## 4. Results

### 4.1 Learning curves

The training converges smoothly. By the end of training the validation accuracy stabilizes around **0.827** with validation loss near **0.438**, while training accuracy approaches **≈0.895**. This gap indicates healthy generalization without severe overfitting.

## 4.2 Final test performance

Two accuracy numbers appear in the notebook. The Keras `model.evaluate` reports **92.91%** because it includes padded positions. The post-hoc analysis removes padding positions and yields a more faithful **per-residue accuracy of 81.77%** on the test set. The latter should be taken as the primary figure.

**Classification report (per residue, padding removed):**

Class	Precision	Recall	F1-score	Support
<b>H</b>	0.77	0.78	0.77	509
<b>C</b>	0.86	0.88	0.87	828
<b>B</b>	0.67	0.52	0.59	106
<b>Overall</b>				<b>1443 residues</b>
<b>Accuracy</b>			<b>0.8177</b>	
<b>Macro avg</b>	0.77	0.73	0.74	
<b>Weighted avg</b>	0.81	0.82	0.82	

**Confusion matrix:** The plot in the notebook shows most errors are between **H** ↔ **C** and **B** ↔ **C**, which is expected given that coil is the dominant and morphologically diverse class and  $\beta$ -strands are comparatively under-represented.

## 4.3 Weight-update GIFs

The generated GIF shows early epochs with higher-magnitude, less structured filters that gradually stabilize as the network converges, consistent with the plateauing of validation metrics after ~80 epochs.

## 5. Discussion

The model achieves **81.8% per-residue Q3 accuracy** on held-out data using only one-hot encodings. Precision and recall are highest for the **coil** class because it dominates the dataset and is easier to fit with local features.  $\beta$ -strand recall is the lowest (**0.52**), reflecting class imbalance and the fact that sheets often require longer-range dependencies than 3-wide convolutions capture. The gap between Keras accuracy and the masked per-residue accuracy highlights the importance of excluding padding from metrics; once masked, the results align with typical Q3 baselines for simple CNNs trained from sequence only.

**Failure modes.** Misclassifications cluster at boundaries between structural states and in long  $\beta$ -strands where long-range context matters. When sequences are much shorter than the global `max_len`, zero-padding can subtly affect batch statistics and make optimization easier on

padded windows; this is mitigated in analysis by masking but still affects training loss and the Keras accuracy.

## 6. Limitations

The model does not use evolutionary profiles or pretrained protein language models, both of which are known to improve secondary-structure prediction. The receptive field is limited to local windows of size three, which constrains the ability to model long-range  $\beta$ -sheet pairings. Padding is not explicitly masked during loss computation, which inflates training and evaluation accuracy unless handled post-hoc.

## 7. Potential Improvements

Long-range context and richer residue features are the main levers. Replacing raw one-hot with either position-specific scoring matrices (PSSMs) or embeddings from pretrained transformers like ESM2 would give the model evolutionary and contextual signals. Architecturally, dilated convolutions, residual blocks, or a lightweight Transformer encoder would increase the receptive field without losing per-position alignment. Introducing a `Masking` layer or per-residue sample weights in Keras would ensure padding is ignored in the loss and metrics. Class imbalance can be addressed with class-weighted loss or focal loss, which should especially help  $\beta$ -strand recall. Finally, hyperparameter tuning for kernel sizes and dropout, early stopping on validation loss, and ensembling can yield incremental gains. Reporting Q8 metrics and segment-overlap measures such as SOV would strengthen the evaluation.

## 8. Reproducibility and File Map

- **Input:** `sequence.fasta`, `Protein Structures/*.ss2`
- **Notebook:** `protein_2d_structure_prediction.ipynb`
- **Artifacts:**
  - Learning curves and confusion matrix are generated inline in the notebook figures.
  - Weight-update GIF: `weights_frames/weights_evolution.gif`

Set a fixed random seed and save figures to files if preparing a static report. Make sure every figure and the GIF are committed to the repository for grading.

## 9. Conclusion

A compact CNN trained on one-hot encodings achieves **81.8% per-residue Q3 accuracy** after masking padding and produces interpretable learning dynamics via weight-evolution GIFs. While already competitive for a minimal setup, the path to higher accuracy is clear: add contextual residue embeddings or profiles, extend the receptive field, handle padding within the training loss, and rebalance the classes. These steps would primarily target  $\beta$ -strand recall and the H $\leftrightarrow$ C boundary, where most residual errors occur.